

Speech Command Recognition

Presentation by
P Venkatesh
AI20MTECH01004



September 26, 2021

Problem Statement

Model Architecture

Dataset Preparation

Experiments and Results

Conclusion

- ▶ The goal of the project is to develop a speech command recognition model
- ▶ keras's neural attention network is used as a baseline framework
- ▶ The model is trained and tested on speech command dataset to recognize 5 different commands as follows:
 - back
 - forward
 - left
 - right
 - stop

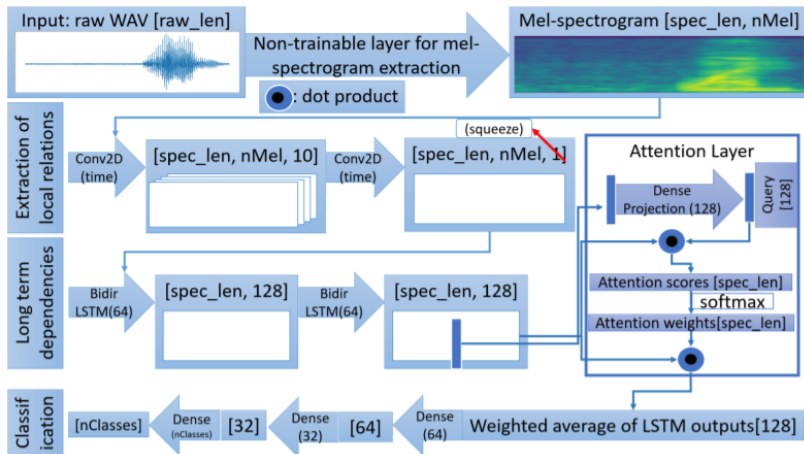



Figure 2: Attention Model



(a) Dataset Preparation using Audacity

- ▶ Set 16KHz as sampling rate
- ▶ Record 80 utterances of each command
- ▶ Save samples of each command in different folders
 - Data/back
 - Data/forward
 - Data/left
 - Data/right
 - Data/stop

Experiments and Results

Model: "Attention"

Layer (type)	Output Shape	Param #	Connected to
Input (InputLayer)	[(None, 49, 39, 1)]	0	
Conv1 (Conv2D)	(None, 49, 39, 10)	60	Input[0][0]
BN1 (BatchNormalization)	(None, 49, 39, 10)	40	Conv1[0][0]
Conv2 (Conv2D)	(None, 49, 39, 1)	51	BN1[0][0]
BN2 (BatchNormalization)	(None, 49, 39, 1)	4	Conv2[0][0]
Squeeze (Reshape)	(None, 49, 39)	0	BN2[0][0]
LSTM_Sequences (LSTM)	(None, 49, 64)	26624	Squeeze[0][0]
FinalSequence (Lambda)	(None, 64)	0	LSTM_Sequences[0][0]
UnitImportance (Dense)	(None, 64)	4160	FinalSequence[0][0]
AttentionScores (Dot)	(None, 49)	0	UnitImportance[0][0] LSTM_Sequences[0][0]
AttentionSoftmax (Softmax)	(None, 49)	0	AttentionScores[0][0]
AttentionVector (Dot)	(None, 64)	0	AttentionSoftmax[0][0] LSTM_Sequences[0][0]
FC (Dense)	(None, 32)	2080	AttentionVector[0][0]
Output (Dense)	(None, 5)	165	FC[0][0]
Total params: 33,184			
Trainable params: 33,162			
Non-trainable params: 22			

Figure 4: Neural Attention Network Architecture

Experiments and Results (cont.)

```
Epoch 1/10  
510/510 - 15s - loss: 0.6114 - sparse_categorical_accuracy: 0.7829 - val_loss: 0.2484 - val_sparse_categorical_accuracy: 0.9112  
Epoch 2/10  
510/510 - 10s - loss: 0.0569 - sparse_categorical_accuracy: 0.9849 - val_loss: 0.0898 - val_sparse_categorical_accuracy: 0.9618  
Epoch 3/10  
510/510 - 10s - loss: 0.0171 - sparse_categorical_accuracy: 0.9960 - val_loss: 0.1021 - val_sparse_categorical_accuracy: 0.9647  
Epoch 4/10  
510/510 - 10s - loss: 0.0027 - sparse_categorical_accuracy: 0.9995 - val_loss: 0.1050 - val_sparse_categorical_accuracy: 0.9597  
Epoch 5/10  
510/510 - 10s - loss: 5.4265e-04 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.1070 - val_sparse_categorical_accuracy: 0.9629  
Epoch 6/10  
510/510 - 10s - loss: 2.8190e-04 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.0999 - val_sparse_categorical_accuracy: 0.9659  
Epoch 7/10  
510/510 - 10s - loss: 1.8118e-04 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.0991 - val_sparse_categorical_accuracy: 0.9653  
Epoch 8/10  
510/510 - 10s - loss: 1.2395e-04 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.1049 - val_sparse_categorical_accuracy: 0.9676  
Epoch 9/10  
510/510 - 10s - loss: 8.5585e-05 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.1056 - val_sparse_categorical_accuracy: 0.9700  
Epoch 10/10  
510/510 - 10s - loss: 6.0929e-05 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.1074 - val_sparse_categorical_accuracy: 0.9703  
<keras.callbacks.History at 0x7fc36019abd0>
```

Figure 5: Results

- ▶ The final accuracy of model is 0.97. Even with this high accuracy the test commands on model give false positives. Hence the model overfits the data.
- ▶ We attribute overfitting to the fact that the data is less and each class has similar data making it hard for the model to generalize.

THANK YOU!